

IMPORT SOME LIBRARIES AND READ THE CSV DATASET

```
import pandas as pd
import numpy as np
import scipy as sp
from sklearn import preprocessing

df = pd.read_csv('loan_approval_dataset.csv')

C:\Users\ashwi\AppData\Local\Temp\ipykernel_12364\679893019.py:1:
DeprecationWarning:
Pyarrow will become a required dependency of pandas in the next major
release of pandas (pandas 3.0),
(to allow more performant data types, such as the Arrow string type,
and better interoperability with other libraries)
but was not found to be installed on your system.
If this would cause problems for you,
please provide us feedback at
https://github.com/pandas-dev/pandas/issues/54466

import pandas as pd
```

HEAD OF THE DATASET

```
df.head()

   loan_id  no_of_dependents      education  self_employed
income_anum \
0           1                  2        Graduate          No
9600000
1           2                  0    Not Graduate         Yes
4100000
2           3                  3        Graduate          No
9100000
3           4                  3        Graduate          No
8200000
4           5                  5    Not Graduate         Yes
9800000

   loan_amount  loan_term  cibil_score
residential_assets_value \
0     29900000       12            778          2400000
1     12200000        8            417          2700000
2     29700000       20            506          7100000
3     30700000       8             467          18200000
```

4	24200000	20	382	12400000
	commercial_assets_value	luxury_assets_value		
bank_asset_value \				
0	17600000	22700000	8000000	
1	2200000	8800000	3300000	
2	4500000	33300000	12800000	
3	3300000	23300000	7900000	
4	8200000	29400000	5000000	
	loan_status			
0	Approved			
1	Rejected			
2	Rejected			
3	Rejected			
4	Rejected			

#### TAIL OF THE DATASET

df.tail()				
	loan_id	no_of_dependents	education	self_employed
income_annum \				
4264	4265	5	Graduate	Yes
1000000				
4265	4266	0	Not Graduate	Yes
3300000				
4266	4267	2	Not Graduate	No
6500000				
4267	4268	1	Not Graduate	No
4100000				
4268	4269	1	Graduate	No
9200000				
	loan_amount	loan_term	cibil_score	
residential_assets_value \				
4264	2300000	12	317	
2800000				
4265	11300000	20	559	
4200000				
4266	23900000	18	457	
1200000				
4267	12800000	8	780	

```

8200000
4268      29700000          10          607
17800000

    commercial_assets_value    luxury_assets_value
bank_asset_value \
4264                  500000          3300000
800000
4265                  2900000          11000000
1900000
4266                  12400000          18100000
7300000
4267                  700000          14100000
5800000
4268                  11800000          35700000
12000000

    loan_status
4264    Rejected
4265    Approved
4266    Rejected
4267    Approved
4268    Approved

```

## EXPLORATORY DATA ANALYSIS(EDA)

### SHAPE OF THE DATASET

```

print("Rows and Columns of the Dataset : ",df.shape)

Rows and Columns of the Dataset : (4269, 13)

```

### INFO OF THE DATASET

provides the essential details about your dataset, such as the number of rows and columns, the number of non-null values, what type of data is in each column, and how much memory your DataFrame is using.

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4269 entries, 0 to 4268
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   loan_id          4269 non-null   int64  
 1   no_of_dependents 4269 non-null   int64  
 2   education        4269 non-null   object  
 3   self_employed    4269 non-null   object  
 4   income_annum    4269 non-null   int64  

```

```

5   loan_amount           4269 non-null    int64
6   loan_term             4269 non-null    int64
7   cibil_score            4269 non-null    int64
8   residential_assets_value 4269 non-null    int64
9   commercial_assets_value 4269 non-null    int64
10  luxury_assets_value    4269 non-null    int64
11  bank_asset_value       4269 non-null    int64
12  loan_status            4269 non-null    object
dtypes: int64(10), object(3)
memory usage: 433.7+ KB

```

## DESCRIPTION OF THE DATASET

The pandas.describe function is used to get a descriptive statistics summary of a given dataframe. This includes mean, count, std deviation, percentiles, and min-max values of all the features.

```

df.describe()

      loan_id  no_of_dependents  income_annum  loan_amount \
count  4269.000000        4269.000000  4.269000e+03  4.269000e+03
mean   2135.000000          2.498712  5.059124e+06  1.513345e+07
std    1232.498479          1.695910  2.806840e+06  9.043363e+06
min    1.000000            0.000000  2.000000e+05  3.000000e+05
25%   1068.000000            1.000000  2.700000e+06  7.700000e+06
50%   2135.000000            3.000000  5.100000e+06  1.450000e+07
75%   3202.000000            4.000000  7.500000e+06  2.150000e+07
max   4269.000000            5.000000  9.900000e+06  3.950000e+07

      loan_term  cibil_score  residential_assets_value \
count  4269.000000        4269.000000  4.269000e+03
mean   10.900445         599.936051  7.472617e+06
std    5.709187          172.430401  6.503637e+06
min    2.000000          300.000000  -1.000000e+05
25%   6.000000          453.000000  2.200000e+06
50%   10.000000          600.000000  5.600000e+06
75%   16.000000          748.000000  1.130000e+07
max   20.000000          900.000000  2.910000e+07

      commercial_assets_value  luxury_assets_value
bank_asset_value
count                  4.269000e+03          4.269000e+03
4.269000e+03
mean                  4.973155e+06          1.512631e+07
4.976692e+06
std                   4.388966e+06          9.103754e+06
3.250185e+06
min                  0.000000e+00          3.000000e+05
0.000000e+00
25%                  1.300000e+06          7.500000e+06

```

```

2.300000e+06
50%           3.700000e+06           1.460000e+07
4.600000e+06
75%           7.600000e+06           2.170000e+07
7.100000e+06
max           1.940000e+07           3.920000e+07
1.470000e+07

```

`df.isnull().sum()`

`.isnull().sum()` returns a DataFrame where each cell is either True or False depending on that cell's null status.

To count the number of nulls in each column we use an aggregate function for summing:

```

df.isnull()

      loan_id  no_of_dependents  education  self_employed
income_annum \
0        False          False     False     False
False
1        False          False     False     False
False
2        False          False     False     False
False
3        False          False     False     False
False
4        False          False     False     False
False
...
...
4264    False          False     False     False
False
4265    False          False     False     False
False
4266    False          False     False     False
False
4267    False          False     False     False
False
4268    False          False     False     False
False

      loan_amount  loan_term  cibil_score
residential_assets_value \
0            False     False     False
False
1            False     False     False
False
2            False     False     False
False

```

```
3      False  False  False
False
4      False  False  False
False
...
.
4264     False  False  False
False
4265     False  False  False
False
4266     False  False  False
False
4267     False  False  False
False
4268     False  False  False
False

    commercial_assets_value  luxury_assets_value
bank_asset_value \
0                  False  False
False
1                  False  False
False
2                  False  False
False
3                  False  False
False
4                  False  False
False
...
.
4264                  False  False
False
4265                  False  False
False
4266                  False  False
False
4267                  False  False
False
4268                  False  False
False

    loan_status
0      False
1      False
2      False
3      False
4      False
...
```

```
4264           False
4265           False
4266           False
4267           False
4268           False

[4269 rows x 13 columns]

df.isnull().sum()

loan_id          0
no_of_dependents 0
education        0
self_employed    0
income_annum     0
loan_amount       0
loan_term         0
cibil_score       0
residential_assets_value 0
commercial_assets_value 0
luxury_assets_value 0
bank_asset_value   0
loan_status        0
dtype: int64
```

## ANALYSIS OF EACH COLUMN

if dropna = False it returns the count of null values and if its True then it returns non-null values

1	2	0	Not Graduate	Yes
4100000				
2	3	3	Graduate	No
9100000				
3	4	3	Graduate	No
8200000				
4	5	5	Not Graduate	Yes
9800000				
...	...	...	...	...
...				
4264	4265	5	Graduate	Yes
1000000				
4265	4266	0	Not Graduate	Yes
3300000				
4266	4267	2	Not Graduate	No
6500000				
4267	4268	1	Not Graduate	No
4100000				
4268	4269	1	Graduate	No
9200000				
loan_amount   loan_term   cibil_score				
residential_assets_value \				
0	29900000	12	778	
2400000				
1	12200000	8	417	
2700000				
2	29700000	20	506	
7100000				
3	30700000	8	467	
18200000				
4	24200000	20	382	
12400000				
...	...	...	...	...
...				
4264	2300000	12	317	
2800000				
4265	11300000	20	559	
4200000				
4266	23900000	18	457	
1200000				
4267	12800000	8	780	
8200000				
4268	29700000	10	607	
17800000				
commercial_assets_value   luxury_assets_value				
bank_asset_value \				
0	17600000		22700000	

```

8000000
1           2200000          8800000
3300000
2           4500000          33300000
12800000
3           3300000          23300000
7900000
4           8200000          29400000
5000000
...
.
4264         500000          3300000
800000
4265         2900000          11000000
1900000
4266         12400000         18100000
7300000
4267         700000          14100000
5800000
4268         11800000         35700000
12000000

      loan_status
0      Approved
1      Rejected
2      Rejected
3      Rejected
4      Rejected
...
4264    Rejected
4265    Approved
4266    Rejected
4267    Approved
4268    Approved

[4269 rows x 13 columns]

```

## NORMALIZATION

```

arr = np.array(df['loan_amount'])
#print(arr)
normalized_arr = preprocessing.normalize([arr])
print(normalized_arr)

[[0.02595843 0.01059173 0.02578479 ... 0.02074938 0.01111264
0.02578479]]

from sklearn.preprocessing import MinMaxScaler

normalization_scaling = MinMaxScaler()

```

```

normalized_data = normalization_scaling.fit_transform(df[['cibil_score']])
print(normalized_data)

[[0.79666667]
 [0.195]
 [0.34333333]
 ...
 [0.26166667]
 [0.8]
 [0.51166667]]

```

Split the data into training, testing and validation sets

```

d_col = ['loan_id', 'education', 'self_employed']

df.drop(d_col, axis=1, inplace=True)

df.head()

no_of_dependents    income_annum    loan_amount    loan_term    cibil_score
\0                   2              9600000     29900000      12          778
1                   0              4100000     12200000      8           417
2                   3              9100000     29700000     20          506
3                   3              8200000     30700000      8           467
4                   5              9800000     24200000     20          382

residential_assets_value    commercial_assets_value
luxury_assets_value \
0                           2400000                  17600000
22700000
1                           2700000                  2200000
8800000
2                           7100000                  4500000
33300000
3                           18200000                 3300000
23300000
4                           12400000                 8200000
29400000

bank_asset_value    loan_status
0                 8000000      Approved
1                 3300000      Rejected

```

```

2           12800000    Rejected
3           7900000    Rejected
4           5000000    Rejected

from sklearn.model_selection import train_test_split
y = df['loan_status']
X = df.drop('loan_status', axis = 1)

X_train,X_test,y_train,y_test =
train_test_split(X,y,train_size=0.7,random_state=50)

X_train.info()

<class 'pandas.core.frame.DataFrame'>
Index: 2988 entries, 893 to 1931
Data columns (total 9 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   no_of_dependents    2988 non-null   int64  
 1   income_annum        2988 non-null   int64  
 2   loan_amount          2988 non-null   int64  
 3   loan_term            2988 non-null   int64  
 4   cibil_score          2988 non-null   int64  
 5   residential_assets_value 2988 non-null   int64  
 6   commercial_assets_value 2988 non-null   int64  
 7   luxury_assets_value   2988 non-null   int64  
 8   bank_asset_value     2988 non-null   int64  
dtypes: int64(9)
memory usage: 233.4 KB

```

## STANDARDIZATION

```

from sklearn.preprocessing import StandardScaler

scaling = StandardScaler()

num_col = X_train.select_dtypes(include=['int64','float64']).columns
num_col
X_train[num_col] = scaling.fit_transform(X_train[num_col])

X_train

      no_of_dependents  income_annum  loan_amount  loan_term
cibil_score \
893           1.455023      0.886012      1.061663   -0.506808
0.842310
3749          -1.486923      1.028895      1.239220    1.235818
1.429435
1958           0.866633      -0.292771      -0.048062   -0.855333
1.599198

```

102	-0.310145	-0.257050	-0.236715	1.584343
0.109858				
974	-1.486923	-0.007005	0.373634	0.190242
0.888815				
...	...	...	...	...
...				
3330	-0.898534	0.457364	0.240467	0.190242
0.667917				
70	0.866633	-1.614437	-1.501803	1.584343
1.506189				-
132	0.278244	-1.150068	-1.268760	1.235818
0.779549				-
2014	0.866633	1.350381	2.127000	0.887293
1.423622				
1931	0.278244	0.457364	0.218272	0.887293
1.227159				-
residential_assets_value    commercial_assets_value				
luxury_assets_value \				
893	-0.896438		-0.307055	
0.188991				
3749	1.049592		2.261685	
1.370010				
1958	-0.170538		0.225207	
0.362888				-
102	0.617141		-0.307055	
0.396000				-
974	1.111371		-0.214488	
0.495339				-
...	...		...	
...				
3330	0.416360		-0.168204	
1.071996				
70	-1.128108		-1.093876	
1.444569				-
132	-1.050884		-1.024450	
1.058254				-
2014	1.744603		1.405438	
0.630493				
1931	0.663475		-0.654182	
0.122765				
bank_asset_value				
893	1.338644			
3749	0.415647			
1958	-0.599649			
102	-0.784248			
974	0.477180			
...	...			

```

3330      1.092511
70       -1.338046
132      -0.722715
2014      1.769375
1931     -0.538116

[2988 rows x 9 columns]

X_train.isnull().sum()

no_of_dependents      0
income_annum          0
loan_amount            0
loan_term              0
cibil_score            0
residential_assets_value  0
commercial_assets_value 0
luxury_assets_value    0
bank_asset_value        0
dtype: int64

```

## FEATURE SELECTION

```

from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE

logreg= LogisticRegression()

rfe = RFE(estimator=logreg, n_features_to_select=5) # running RFE with
15 variables as output15

rfe = rfe.fit(X_train,y_train)

col = X_train.columns[rfe.support_]

col

Index(['income_annum', 'loan_amount', 'loan_term', 'cibil_score',
       'luxury_assets_value'],
      dtype='object')

# Assuming you have already fitted the RFE
feature_ranking = rfe.ranking_
selected_features_mask = rfe.support_

# Print the ranking of features
print("Feature Ranking:")
print(feature_ranking)

# Print the selected features based on the boolean mask
selected_features = X_train.columns[selected_features_mask]

```

```
print("\nSelected Features:")
print(selected_features)

Feature Ranking:
[5 1 1 1 1 3 4 1 2]

Selected Features:
Index(['income_annum', 'loan_amount', 'loan_term', 'cibil_score',
       'luxury_assets_value'],
      dtype='object')

logreg_ranking = rfe.estimator_.coef_[0]
print("\nLogistic Regression Coefficient Ranking:")
print(logreg_ranking)

Logistic Regression Coefficient Ranking:
[ 1.31776459 -1.15086366  0.87859794 -4.1817298  -0.22225772]
```